

# Auto-generating Quantized DNN Kernels using TVM - Mila Lukic

The goal of this project is to use TVM compile a quantized ResNet50 down to a executable function; for simplicity, it's not yet necessary to test the accuracy of the compiled model.

This is the github link to the implementation:

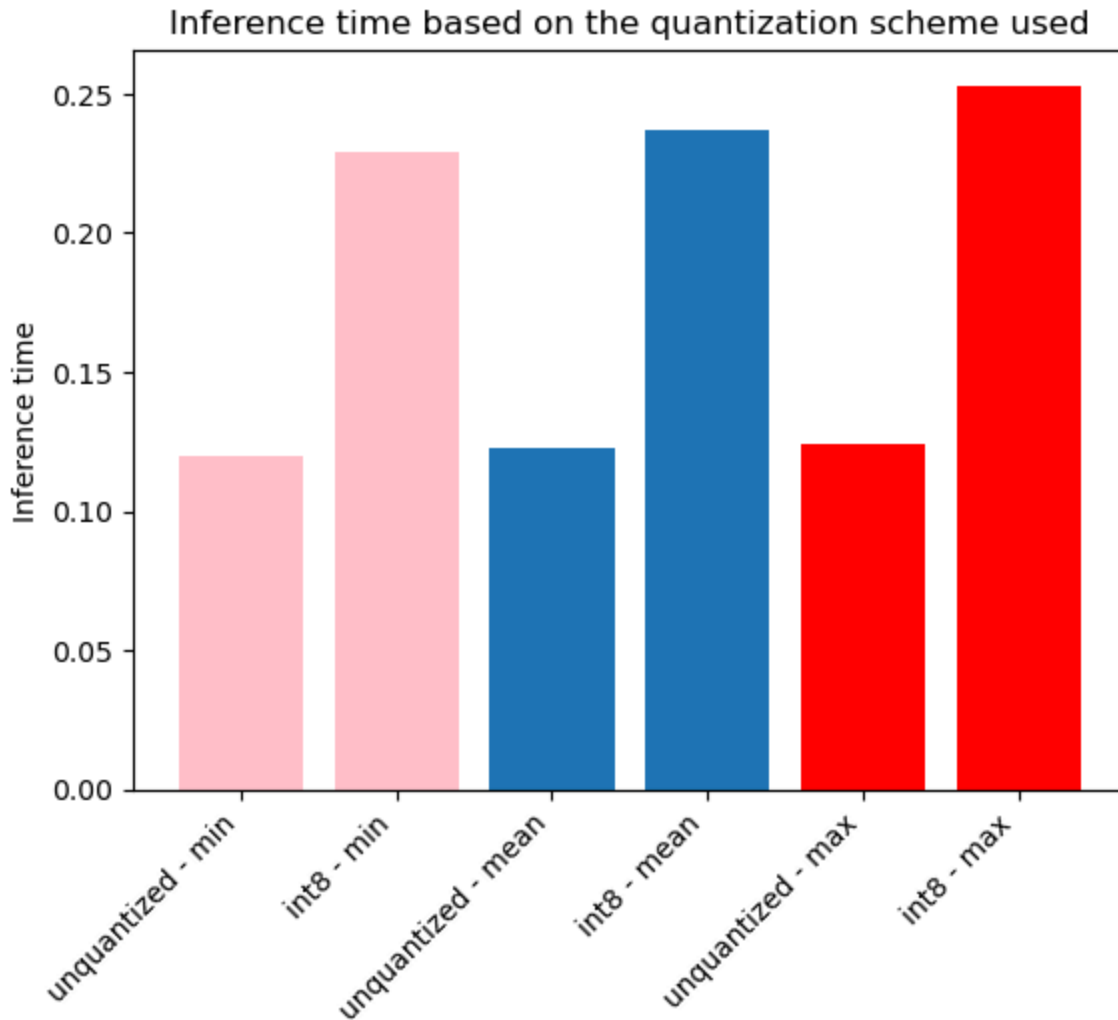
[https://github.com/milalukic/tvm-quantization/blob/main/tvm\\_quantization\\_mila.ipynb](https://github.com/milalukic/tvm-quantization/blob/main/tvm_quantization_mila.ipynb)

In this document, I am going to answer the “non-coding” questions posed in the assignment.

## Model execution time comparison

*Use TVM's utility functions to benchmark the inference time of the quantized model vs. the un-quantized model.*

*In this task we will not try to maximize the performance of the quantized DNN, but if there is no speedup, you should try to understand it and formulate a guess.*



We can see that the unquantized DNN is faster in terms of inference time (when we look at all of the benchmarks - min, max, as well as mean). This is unexpected - the quantization to a smaller precision is supposed to make the DNN faster, but less accurate.

After some research, I found out these possible reasons for this model to be slower:

- **Overhead from Quantization:** While the lower precision reduces computation for the core operations, these calibration and conversion might outweigh the gains on some hardware.
- **Hardware Acceleration Not Optimized:** If my hardware is not specifically optimized for running low-precision models, it might not be able to fully exploit the benefits of quantization.

Considering the fact that I've got this warning message when I compiled the unquantized DNN:

One or more operators have not been tuned. Please tune your model for better performance. Use DEBUG logging level to see more details.

I looked a bit more into it - it wasn't intuitive at all to me that the model that's supposed to be "tuned for better performance" is still performing better compared to the model that didn't get the same message!

I researched it online, but couldn't find a specific solution that made sense. As a last resort, I decided to ask Gemini (Google's chatbot) and among other things that it mentioned, I came across this paragraph that I found interesting:

*It's possible that the lack of tuning in the unquantized model leads to a simpler execution path. Without optimizations, the model might be using a more generic implementation that, in some cases, could be surprisingly efficient for the specific hardware you're running on.*

This, paired with possible overhead and hardware issues, made sense to me.

## How did TVM know that I wanted to quantize to int8?

*In your quantization setup, how did TVM know that you wanted to quantize to int8? Look into that, and vary the number of bits of quantization (the in int-). Searching in forum and peeking the source code of the quantizer class will both help.*

The `qconfig` function allows us to define the quantization configuration using the `QConfig` class. This class has attributes like `dtype_input` and `dtype_weight`, which specify the number of bits and data type for different elements (input, weight) in the model.

When quantizing the DNN in my code, I specifically set the configuration so that `dtype_input` and `dtype_weight` were equal to "int8".

```
with relay.quantize.qconfig(calibrate_mode='global_scale',
                           dtype_input="int8",
                           dtype_weight="int8",
                           dtype_activation="int8",
                           global_scale=8.0):
    mod = relay.quantize.quantize(mod, params)
```

Upon taking a closer look at the source code of the `quantize` class (<https://github.com/apache/tvm/blob/main/python/tvm/relay/quantize/quantize.py>), I came across some more interesting discoveries.

While `qconfig` allows us to specify the number of bits, it also provides defaults. By default, `nbit_input` and `nbit_weight` are both set to 8, corresponding to int8. `dtype_input` and `dtype_weight` are also set to "int8" by default.

So to sum up: while we can influence the quantization type through qconfig, TVM's defaults and calibration process are geared towards using int8 for efficient low-precision inference. That's how TVM will know that we want to quantize to "int8", even if we don't specifically set the configuration to that.

*Try out int8 -> int4 -> int2 -> int1; note which precisions work. When it doesn't work, note exactly which part is failing.*

```
In [20]: keys = ["int8", "int4", "int2", "bool"]
for key in keys:
    try:
        lib = f_quantize(key)
        m = graph_executor.GraphModule(lib["default"](dev))
        m.set_input("input0", tvm.nd.array(img.astype(dtype)))
    except Exception as e:
        print(key + " is failing!")
```

```
int4 is failing!
int2 is failing!
```

We can see from the output that `int4` and `int2` are failing. Before setting up the try-catch block we see here, I tried compiling them one by one and the result I received was "TVMError: unknown data type int4" (as seen below). This makes me think that they are probably not supported by the TMV compiler.

```
at /home/mila/Desktop/uiuc/tvm/include/tvm/relay/expr_functor.h:95
6: tvm::relay::ExprFunctor<tvm::RelayExpr (tvm::RelayExpr const&)>::InitVTable()::{lambda(tvm::runtime::ObjectRef const&, tvm::relay::ExprFunctor<tvm::RelayExpr (tvm::RelayExpr const&)>*)#6}::_FUN(tvm::runtime::ObjectRef const&, tvm::relay::ExprFunctor<tvm::RelayExpr (tvm::RelayExpr const&)>*)
at /home/mila/Desktop/uiuc/tvm/include/tvm/relay/expr_functor.h:128
5: tvm::relay::ExprFunctor<tvm::RelayExpr (tvm::RelayExpr const&)>::InitVTable()::{lambda(tvm::runtime::ObjectRef const&, tvm::relay::ExprFunctor<tvm::RelayExpr (tvm::RelayExpr const&)>*)#6}::operator()(tvm::runtime::ObjectRef const&, tvm::relay::ExprFunctor<tvm::RelayExpr (tvm::RelayExpr const&)>*) const
at /home/mila/Desktop/uiuc/tvm/include/tvm/relay/expr_functor.h:128
4: tvm::relay::MixedModeMutator::VisitExpr_(tvm::relay::CallNode const*)
at /home/mila/Desktop/uiuc/tvm/include/tvm/relay/expr_functor.h:291
3: tvm::RelayExpr tvm::relay::MixedModeMutator::Rewrite<tvm::relay::CallNode>(tvm::relay::CallNode const*)
at /home/mila/Desktop/uiuc/tvm/include/tvm/relay/expr_functor.h:313
2: tvm::relay::ForwardRewriter::Rewrite_(tvm::relay::CallNode const*, tvm::RelayExpr const&)
at /home/mila/Desktop/uiuc/tvm/src/relay/transforms/forward_rewrite.cc:155
1: tvm::relay::quantize::QuantizeRealize(tvm::relay::Call const&, tvm::runtime::Array<tvm::RelayExpr, void> const&, tvm::runtime::ObjectRef const&)
at /home/mila/Desktop/uiuc/tvm/src/relay/quantize/realize.cc:126
0: tvm::relay::Constant tvm::relay::MakeConstantScalar<int>(tvm::runtime::DataType, int)
at /home/mila/Desktop/uiuc/tvm/src/relay/quantize/./qnn/./op/././op/./transforms/pattern_utils.h:290
File "/home/mila/Desktop/uiuc/tvm/src/relay/quantize/./qnn/./op/././op/./transforms/pattern_utils.h", line 2
90
TVMError: unknown data type int4
```